

Search

Dra. M^a Dolores Rodríguez Moreno

Objectives

Specific Objectives

- To understand the role of the search in AI
- Main search algorithms

Source

- Stuart Russell & Peter Norvig (2009). Artificial Intelligence: A Modern Approach. (3rd Edition). Ed. Pearsons.

Outline

- **Introduction**
- Problem formulation
- Problem types
- Basic search algorithms
- Conclusions

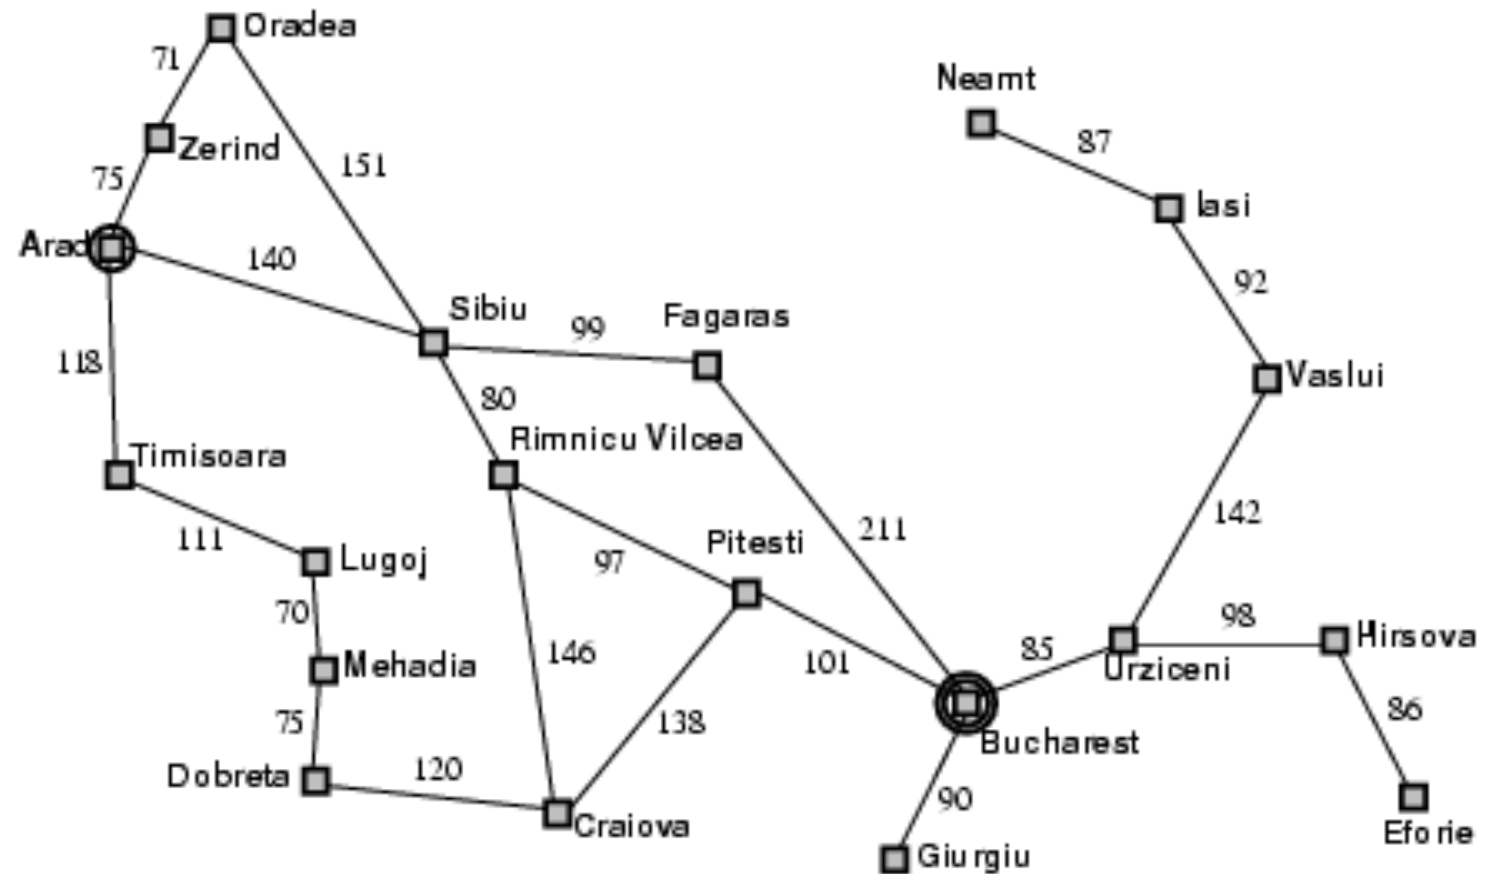
Introduction

- Early AI works were directed to:
 - Proof of theorems
 - Solving crosswords
 - Games
- All in AI is search
 - Not entirely true (obviously) but more than we can imagine
 - Finding a good/best solution to a problem among several possible solutions

Outline

- Introduction
- **Problem formulation**
- Problem types
- Basic search algorithms
- Conclusions

Problem formulation



Problem formulation (I)

- Agent must maximize its performance measure
- Example: On holiday in Romania; currently in Arad
Flight leaves tomorrow from Bucharest
- Formulate goal:
 - be in Bucharest
- Formulating the problem:
 - states: multiple cities
 - actions: drive between cities
- Finding a solution:
 - Sequence cities, eg., Arad, Sibiu, Fagaras, Bucharest
- Process of finding such a solution: **Search**

Problem formulation (II)

- Assumptions of the environment:
 - Static: search and formulation is done without considering changes in the environment
 - Observable: the initial state is known
 - Discrete: the alternative locations are known
 - Deterministic: each state is determined by the current state and the action executed
- The solutions are simple sequences of actions, they are executed without considering perceptions

Problem formulation (III)

- A problem is defined by four items:
 1. initial state e.g., "at Arad"
 2. actions or successor function $S(x)$ = set of action–state pairs
e.g., $S(\text{Arad}) = \{ \langle \text{Arad} \rightarrow \text{Zerind}, \text{Zerind} \rangle, \dots \}$
 3. goal test, can be
 - explicit, e.g., $x = \text{"at Bucharest"}$
 - implicit, e.g., $\text{Checkmate}(x)$
 4. path cost (additive)
 - e.g., sum of distances, number of actions executed, etc.
 - $c(x,a,y)$ is the step cost, assumed to be ≥ 0
- A solution is a sequence of actions leading from the initial state to a goal state

Problem formulation (IV)

- Real world is absurdly complex
 - state space must be abstracted for problem solving
- (Abstract) state = set of real states
- (Abstract) action = complex combination of real actions
 - e.g., "Arad → Zerind" represents a complex set of possible routes, detours, rest stops, etc.
- For guaranteed realizability, any real state "in Arad" must get to some real state "in Zerind"
- (Abstract) solution =
 - set of real paths that are solutions in the real world
- Each abstract action should be "easier" than the original problem

Outline


- Introduction
- Problem formulation
- **Problem types**
- Basic search algorithms
- Conclusions

Problem types

- Deterministic, fully observable → **single-state problem**
 - Agent knows exactly which state it will be in; solution is a sequence
- Non-observable → **sensorless problem** (conformant problem)
 - Agent may have no idea where it is; solution is a sequence
- Nondeterministic and/or partially observable → **contingency problem**
 - percepts provide new information about current state
 - often interleave search with execution
- Unknown state space → **exploration problem**

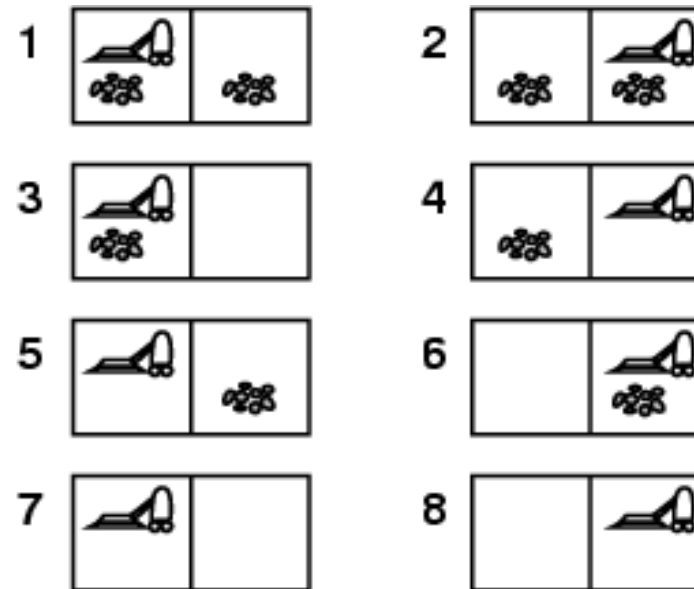
Problem types: example

- Single-state, start in #5.
Solution?



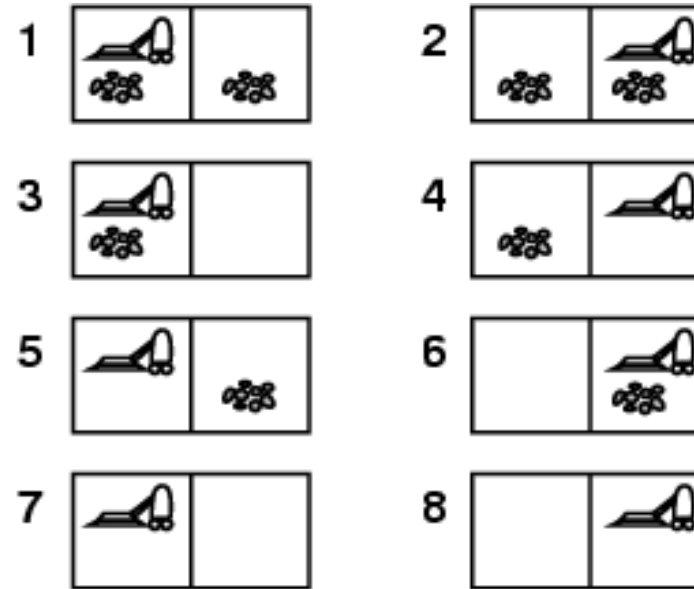
Problem types: example (I)

- Single-state, start in #5.
Solution? [*Right, Suck*]




Problem types: example (II)

- Single-state, start in #5.
Solution? *[Right, Suck]*
- Sensorless, start in $\{1,2,3,4,5,6,7,8\}$ e.g.,
Right goes to $\{2,4,6,8\}$
Solution?




Problem types: example (III)

- Single-state, start in #5.
Solution? *[Right, Suck]*
- Sensorless, start in {1,2,3,4,5,6,7,8} e.g.,
Right goes to {2,4,6,8}
Solution? *[Right, Suck, Left, Suck]*



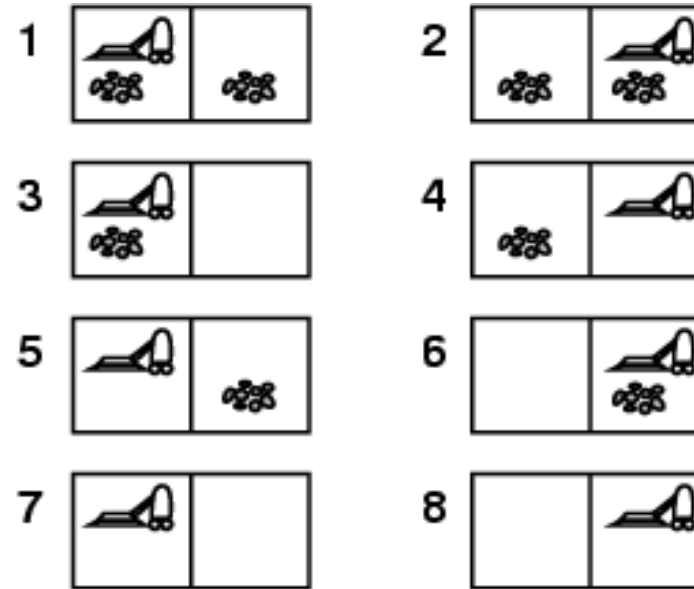
Problem types: example (IV)

- Single-state, start in #5.
Solution? *[Right, Suck]*
- Sensorless, start in {1,2,3,4,5,6,7,8} e.g.,
Right goes to {2,4,6,8}
Solution? *[Right, Suck, Left, Suck]*
- Contingency
 - Nondeterministic: suck may dirt the carpet
 - Partially observable: location, dirt at current location
 - Percept: [L, Clean], i.e., start in #5 or #7**Solution?**



Problem types: example (V)

- Single-state, start in #5.
Solution? *[Right, Suck]*
- Sensorless, start in {1,2,3,4,5,6,7,8} e.g.,
Right goes to {2,4,6,8}
Solution? *[Right, Suck, Left, Suck]*
- Contingency
 - Nondeterministic: Suck may dirt the carpet
 - Partially observable: location, dirt at current location
 - Percept: [L, Clean], i.e., start in #5 or #7
Solution? *[Right, **if** dirt **then** Suck]*



Type of environment

- **Fully observable** (vs. partially observable): an agent's sensors give it access to the complete state of the environment at each point in time
- **Deterministic** (vs. non-deterministic): the next state of the environment is completely determined by the current state and the action executed by the agent
- **Episodic** (vs. sequential): the agent's experience is divided into atomic "episodes" (each episode consists of the agent perceiving and then performing a single action), and the choice of action in each episode depends only on the episode itself
- **Discrete** (vs. continuous): A limited number of distinct, clearly defined percepts and actions
- **Static** (vs. dynamic): if the environment does not change when the agent is deliberating
 - Semidynamic: the environment does not change but the performance of the agent
- **Single agent** (vs. multiagent): an agent operating by itself in an environment

Environment	Observable	Deterministic	Episódic	Static	Discrete	Agents
<i>Chess</i>						
<u>Chess with clock</u>						
<i>Poker</i>						
<u>Taxi</u>						
<u>Medical diagnosis</u>						
<u>Image analysis</u>						
<u>Clasifier</u> <u>Robot</u>						
<u>Refinery controller</u>						
<u>Interactive Tutor</u>						

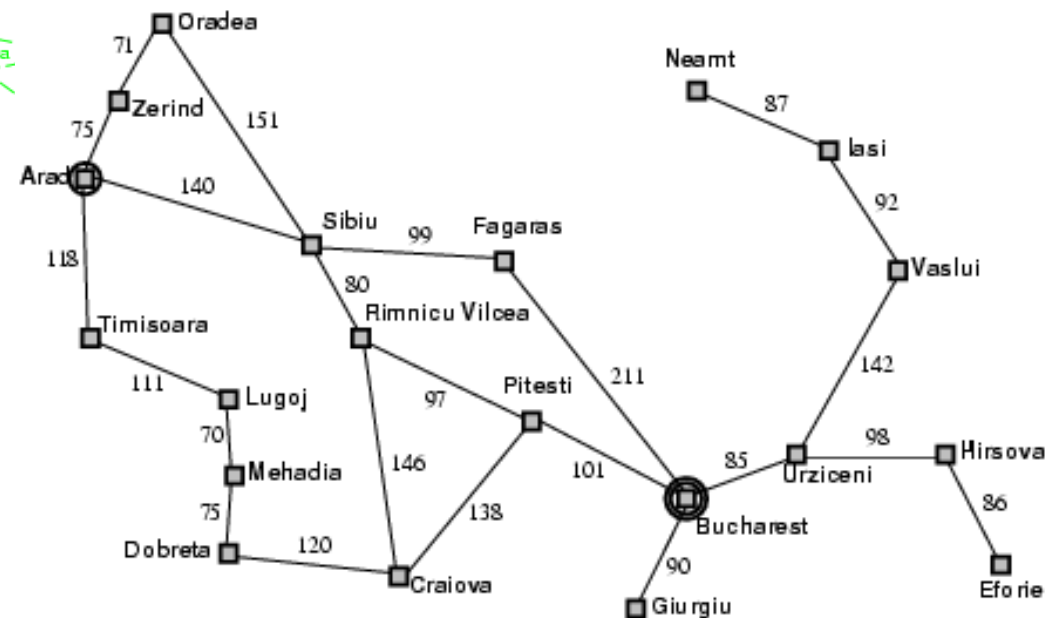

Outline

- Introduction
- Problem formulation
- Problem types
- **Search algorithms**
- Conclusions


Search algorithms

- We have formulated problems, we now need to solve them: search tree
- In general we can have a search graph rather than a tree when the state can be reached from multiple paths
- Basic idea:
 - offline, simulated exploration of state space by generating successors of already-explored states (i.e. **expanding** states)


Search algorithms: tree search example



Search algorithms: tree search example




Search algorithms: tree search example




Search algorithms: states vs. nodes

- A **state** is a (representation of) a physical configuration
- A **node** is a data structure constituting part of a search tree includes **state**, **parent node**, **action**, **path cost $g(x)$** , **depth**




- The Expand function to create the corresponding states

Search algorithms: Vacuum world (I)



- states?
- actions?
- goal test?
- path cost?

Search algorithms: Vacuum world (II)



- states? integer dirt and robot location
- actions? *Left, Right, Suck*
- goal test? no dirt at all locations
- path cost? 1 per action

Search algorithms: The 8-puzzle (I)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- states?
- actions?
- goal test?
- path cost?

Search algorithms: The 8-puzzle (II)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- states? locations of tiles
- actions? move blank left, right, up, down
- goal test? = goal state (given)
- path cost? 1 per move

[Note: optimal solution of n -Puzzle family is NP-hard]

Search strategies

- A search strategy is defined by picking the order of node expansion
- Strategies are evaluated along the following dimensions:
 - completeness: does it always find a solution if one exists?
 - time complexity: number of nodes generated
 - space complexity: maximum number of nodes in memory
 - optimality: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
 - b: maximum branching factor of the search tree
 - d: depth of the least-cost solution
 - m: maximum depth of the state space (may be ∞)

Outline


- Introduction
- Problem formulation
- Problem types
- Basic search algorithms
 - **Uninformed search**
 - Informed search
- Conclusions

Introduction

- Uninformed search strategies use only the information available in the problem definition
 - Breadth-first search/ Búsqueda en anchura
 - Uniform-cost search/ Búsqueda de coste uniforme
 - Depth-first search/ Búsqueda en profundidad
 - Depth-limited search/ Búsqueda en profundidad limitada
 - Iterative deepening search/Búsqueda de profundización iterativa


Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
 - *fringe* is a FIFO queue, i.e., new successors go at end




Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
 - *fringe* is a FIFO queue, i.e., new successors go at end




Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
 - *fringe* is a FIFO queue, i.e., new successors go at end



Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
 - *fringe* is a FIFO queue, i.e., new successors go at end



Properties of breadth-first search

- Complete? Yes (if b is finite)
- Time? $1 + b + b^2 + b^3 + \dots + b^d + b(b^d - 1) = O(b^{d+1})$
- Space? $O(b^{d+1})$ (keeps every node in memory)
- Optimal? Yes (if cost = 1 per step)



Space is the bigger problem (more than time)

Each state has b successors (branching factor)

d is the shallower depth

Uniform-cost search

- Expand least-cost unexpanded node
- **Implementation:**
 - *fringe* = queue ordered by path cost
- Find the solution with minimum cumulative cost, i.e. an optimal solution



Uniform-cost search (Solution)

Initialization: { [S , 0] }

Iteration1: { [S->A , 1] , [S->G , 12] }


Iteration2: { [S->A->C , 2] , [S->A->B , 4] , [S->G , 12] }

Iteration3: { [S->A->C->D , 3] , [S->A->C->G , 4] , [S->A->B->D , 7] , [S->G , 12] }

Iteration 4: { [S->A->C->D->G , 6] , [S->A->C->G , 4] , [S->A->B->D , 7] , [S->G , 12] }

Iteration 5: { [S->A->C->G , 4] , [S->A->C->D->G , 6] , [S->A->B->D->G , 10] , [S->G , 12] }

Solution: S->A->C->G.




Uniform-cost search

- Complete? Yes, if step cost $\geq \epsilon$
- Time? # of nodes with $g \leq$ cost of optimal solution, $O(b^{\text{ceiling}(C^*/\epsilon)})$ where C^* is the cost of the optimal solution
- Space? # of nodes with $g \leq$ cost of optimal solution, $O(b^{\text{ceiling}(C^*/\epsilon)})$
- Optimal? Yes – nodes expanded in increasing order of $g(n)$

If all costs are equal $\rightarrow O(b^d)$


Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front



Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front




Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front



Depth-first search

- Expand deepest unexpanded node
- **Implementation:**
 - *fringe* = LIFO queue, i.e., put successors at front



Properties of depth-first search

- Complete? No: fails in infinite-depth spaces, spaces with loops
 - Modify to avoid repeated states along path
 - complete in finite spaces
- Time? $O(b^m)$: terrible if m is much larger than d
 - but if solutions are dense, may be much faster than breadth-first
- Space? $O(bm)$
- Optimal? No

Depth-limited search

- = depth-first search with depth limit L
- i.e., nodes at depth L have no successor

Iterative deepening search $L=0$

Limit = 0




Iterative deepening search $L=1$

Limit = 1




Iterative deepening search $L=2$



Iterative deepening search $L=3$

Limit = 3





Properties of iterative deepening search

- Complete? Yes
- Time? $(d+1)b^0 + d b^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- Space? $O(bd)$
- Optimal? Yes, if step cost = 1

Repeated states

- Failure to detect repeated states can turn a linear problem into an exponential one!



Summary of algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes	No	No	Yes
Time	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes	Yes	No	No	Yes

Outline

- Introduction
- Problem formulation
- Problem types
- Basic search algorithms
 - Uninformed search
 - **Informed search**
- Conclusions

Informed search strategies

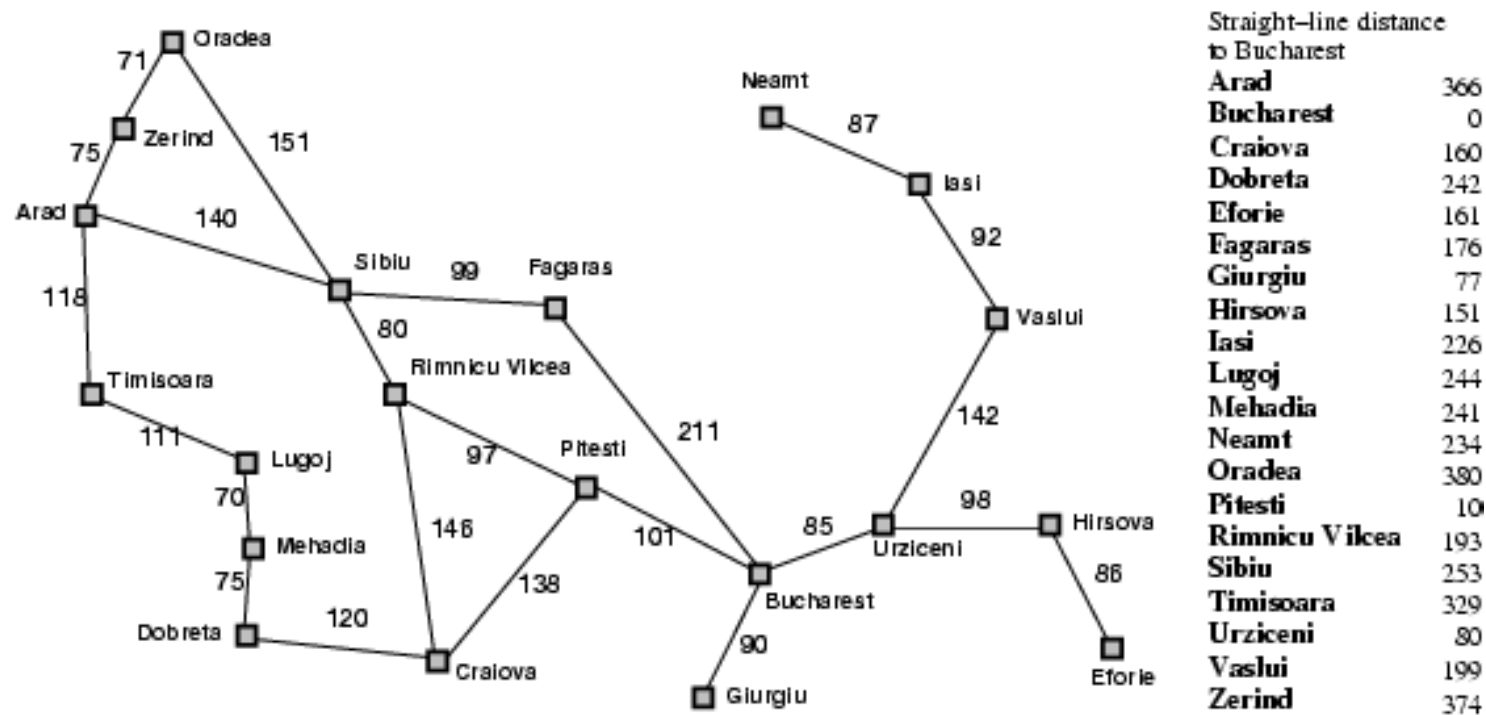
- Use the problem-specific knowledge beyond the definition of the problem itself to find more efficient solutions than uninformed strategies
 - Best-first search
 - Greedy best-first search/Búsqueda voraz primero el mejor
 - A^*
 - Heuristics
 - Local search algorithms
 - Hill-climbing search/Búsqueda de escalada
 - Simulated annealing search/Búsqueda de temple simulado
 - Local beam search/Búsqueda de haz local
 - Genetic algorithms/Algoritmos genéticos

Best-first search

- Idea: use an **evaluation function** $f(n)$ for each node
 - estimate of "desirability"
 - Expand most desirable unexpanded node
- Implementation:

Order the nodes in fringe in decreasing order of desirability
- Special cases:
 - greedy best-first search
 - A* search

Romania with step costs in km



Informed search strategies

- Best-first search
 - Greedy best-first search
 - A* search
- Heuristics
- Local search algorithms
 - Hill-climbing search
 - Simulated annealing search
 - Local beam search
 - Genetic algorithms

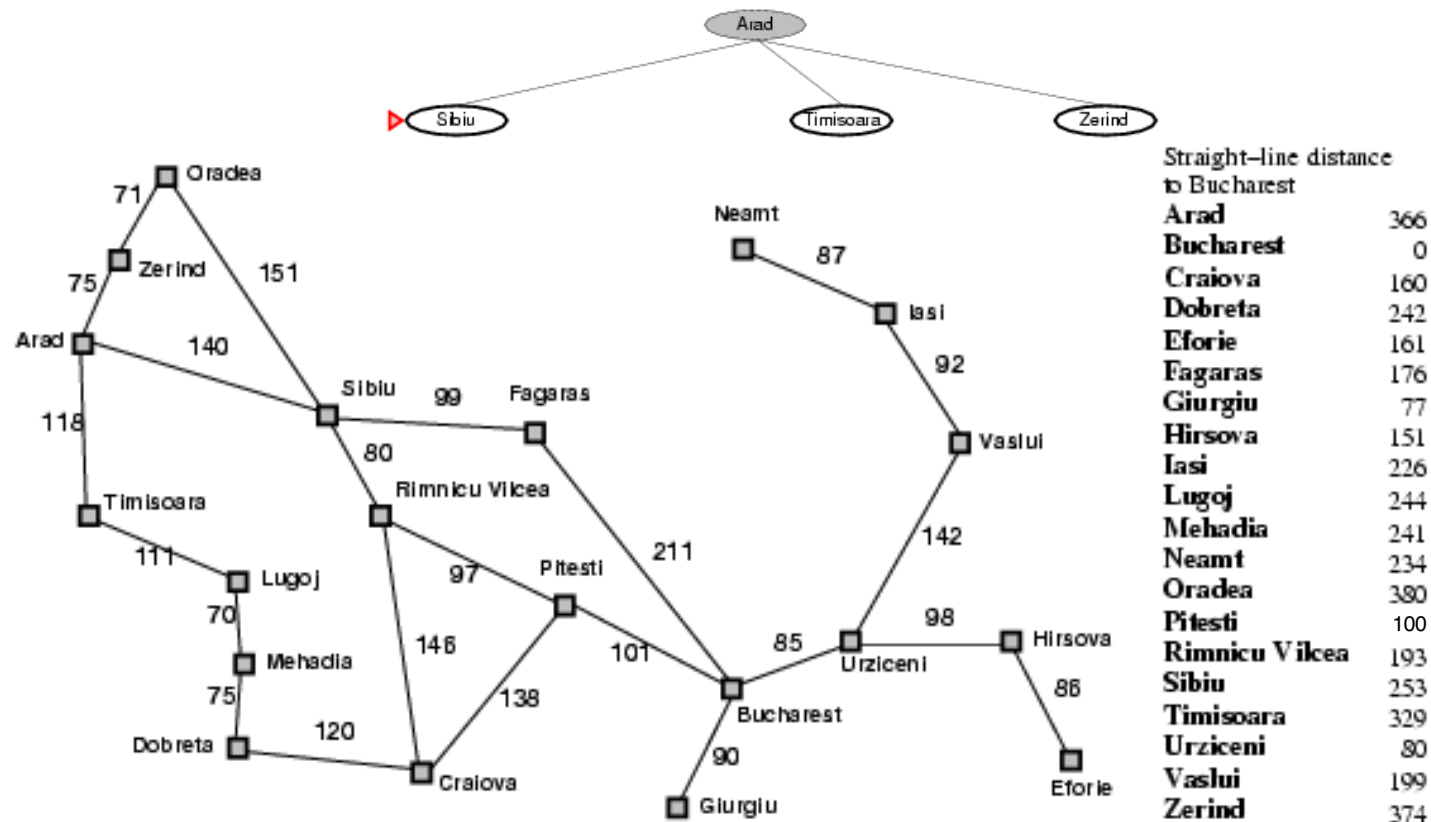
Greedy best-first search

- Evaluation function $f(n) = h(n)$ (**h**euristic) = estimate of cost from n to *goal*
- Greedy best-first search expands the node that **appears** to be closest to the goal
- Implementation: as a priority queue to keep the fringe in ascending order of f -values
- e.g., $h_{SLD}(n)$ = straight-line distance from n to Bucharest


Greedy best-first search example



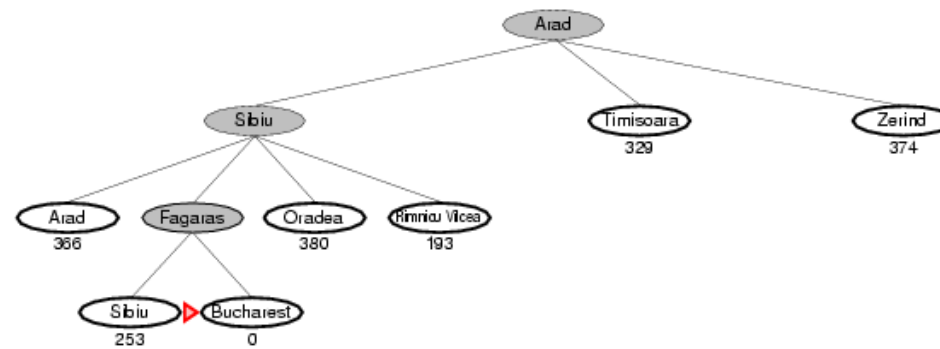
Greedy best-first search example



Greedy best-first search example



Greedy best-first search example



Properties of greedy best-first search

- Complete? No – can get stuck in loops, e.g., Iasi \rightarrow Neamt \rightarrow Iasi \rightarrow Neamt \rightarrow
- Time? $O(b^m)$, but a good heuristic can give dramatic improvement
- Space? $O(b^m)$ -- keeps all nodes in memory
- Optimal? No

Similar to depth-first search

Each state has b successors (branching factor)

d is the depth of the shallowest solution

Informed search strategies

- Best-first search
 - Greedy best-first search
 - A^* search
- Heuristics
- Local search algorithms
 - Hill-climbing search
 - Simulated annealing search
 - Local beam search
 - Genetic algorithms


A* search

- Idea: avoid expanding paths that are already expensive
- Evaluation function $f(n) = g(n) + h(n)$
 - $g(n)$ = cost so far to reach n
 - $h(n)$ = estimated cost from n to goal
 - $f(n)$ = estimated total cost of path through n to goal
- A* is optimal if $h(n)$ is an admissible heuristic such that $h(n)$ never overestimates the cost to reach the goal

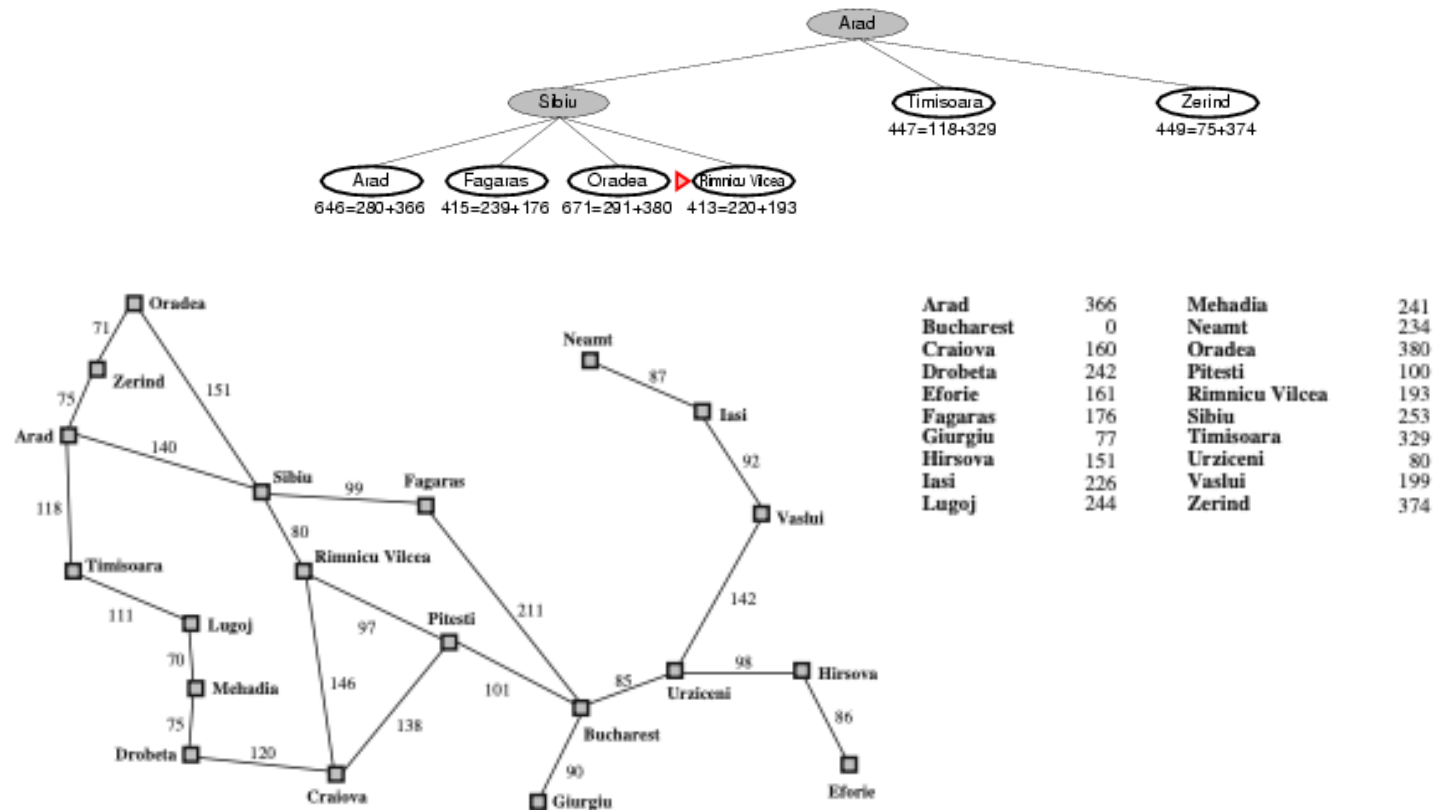
A* search example




A* search example



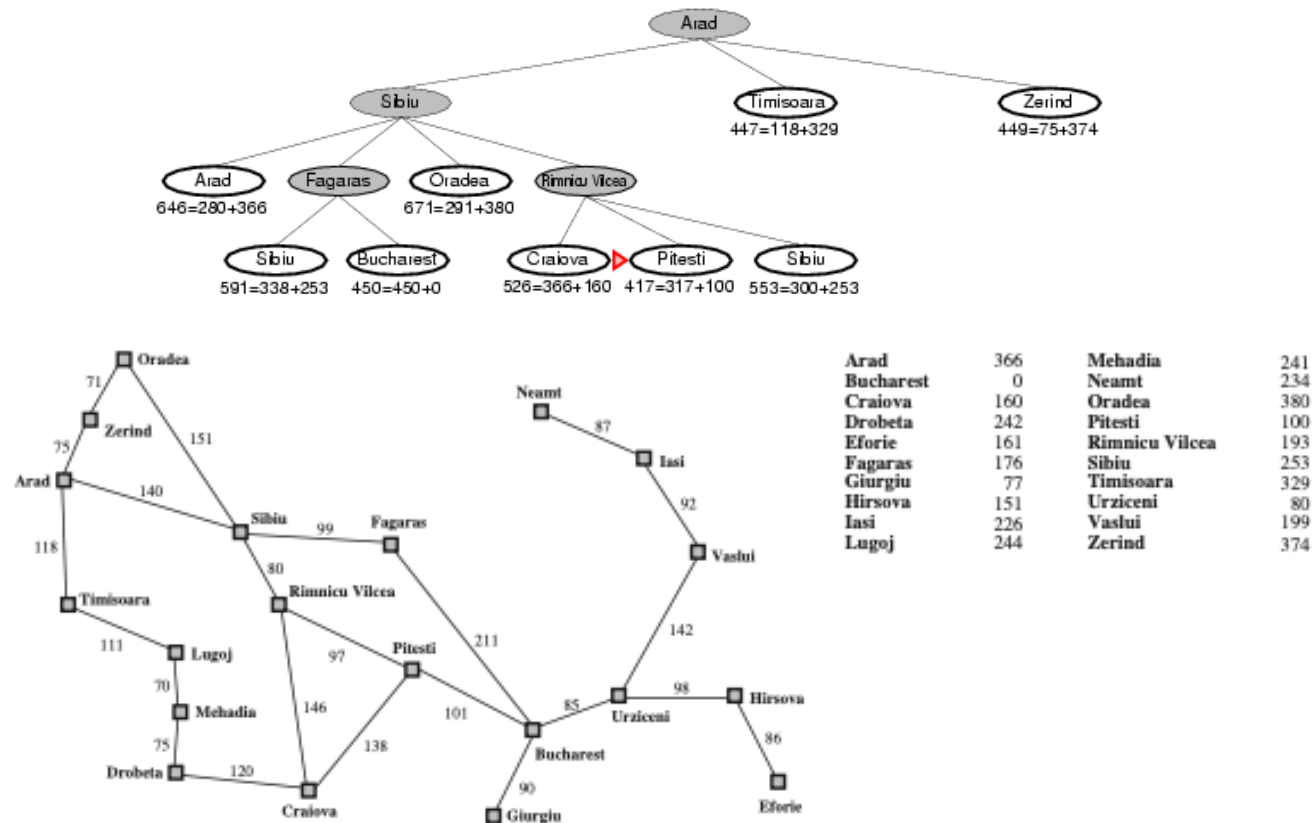
A* search example



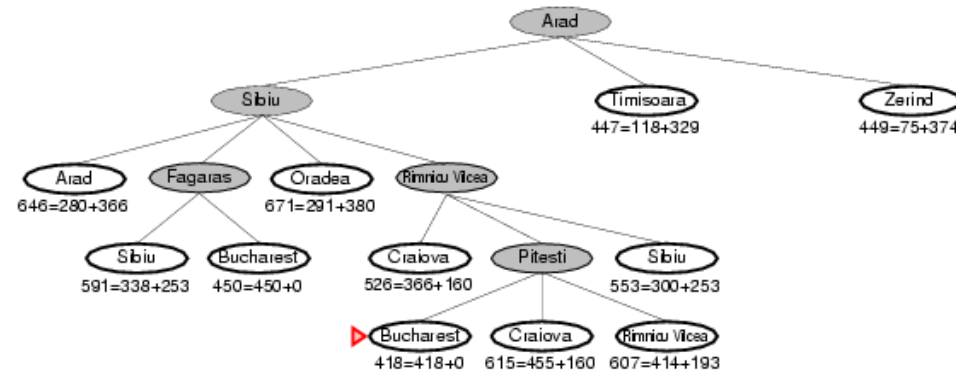
A* search example



A* search example



A* search example



Properties of A*

- Complete? Yes (unless there are infinitely many nodes with $f \leq f(G)$)
- Time? Exponential
- Space? Keeps all nodes in memory
- Optimal? Yes

Informed search strategies

- Best-first search
- Greedy best-first search
- A* search
- **Heuristics**
- Local search algorithms
 - Hill-climbing search
 - Simulated annealing search
 - Local beam search
 - Genetic algorithms

Admissible heuristics

- Objective:
 - produce a solution that is good enough for solving the problem at hand
 - This solution may not be the best but approximate the exact solution
- $h(n)$ is **admissible** if for node n , $h(n) \leq h^*(n)$, where $h^*(n)$ is the **true** cost to reach the goal state from n
- An admissible heuristic **never overestimates** the cost to reach the goal, i.e., it is **optimistic**
- Example: $h_{SLD}(n)$ (never overestimates the actual road distance)
- **Theorem:** If $h(n)$ is admissible, A^* using TREE-SEARCH is optimal

Admissible heuristics for 8-puzzle

- Heuristic: produce a solution in a reasonable time frame, good enough to solve the problem
- The average cost for the 8-puzzle are approx. 22 steps. Here are 26 steps.
- Branching factor is approx. 3
 - Empty in the middle, 4 mov
 - Empty in the corner, 2 mov
 - Rest cases, 3 mov
- Depth-first search will look 3^{22} states
- If we keep track of repeated states, we could reduce to 170.000
- In the 15-puzzle = 10^{13}

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Admissible heuristics for 8-puzzle

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance

(i.e., + horizontal and vertical distance from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = ?$
- $h_2(S) = ?$

Admissible heuristics for 8-puzzle

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance
(i.e., + horizontal and vertical distance from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = ?$ 8
- $h_2(S) = ?$ $3+1+2+2+2+3+3+2 = 18$

Dominance

- If $h_2(n) \geq h_1(n)$ for all n (both admissible) then h_2 **dominates** h_1
- h_2 is better for search
- Typical search costs (average number of nodes expanded):
 - $d=12$ IDS = 3,644,035 nodes
 $A^*(h_1) = 227$ nodes
 $A^*(h_2) = 73$ nodes
 - $d=24$ IDS = too many nodes
 $A^*(h_1) = 39,135$ nodes
 $A^*(h_2) = 1,641$ nodes

Relaxed problems

- A problem with fewer restrictions on the actions is called a **relaxed problem**
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem
- If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then $h_1(n)$ gives the shortest solution
- If the rules are relaxed so that a tile can move to **any adjacent square**, then $h_2(n)$ gives the shortest solution

Informed search strategies

- Best-first search
 - Greedy best-first search
 - A^* search
- Heuristics
- **Local search algorithms**
 - Hill-climbing search
 - Simulated annealing search
 - Local beam search

Local search algorithms

- In many optimization problems, the **path** to the goal is irrelevant; the goal state itself is the solution
- State space = set of "complete" configurations. Find configuration satisfying constraints, e.g., n-queens
- In such cases, we can use **local search algorithms**
- Keep a single "current" state, try to improve it
- Work with one current state and generally moves to the neighboring state
- The paths followed by the search are not retained
 - They use little memory
 - You can find reasonable solutions in large state spaces or infinite

Example: n -queens

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal



Informed search strategies


- Best-first search
- Greedy best-first search
- A* search
- Heuristics
- Local search algorithms
 - **Hill-climbing search**
 - Simulated annealing search
 - Local beam search

Hill-climbing search

- It's just a loop that moves in the direction of increasing value
 - Ends when it reaches a peak where no neighbor has a higher value
 - The search tree is not kept, just a data structure of the current node to check the goal condition and its objective function value
- "Like climbing Everest in thick fog with amnesia"

Hill-climbing search

- Problem: depending on initial state, can get stuck in local maxima




Hill-climbing search: 8-queens problem

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♚	13	16	13	16
♚	14	17	15	♚	14	16	16
17	♚	16	18	15	♚	15	♚
18	14	♚	15	15	14	♚	16
14	14	13	17	12	14	12	18

- h = number of pairs of queens that are attacking each other, either directly or indirectly
- $h = 17$ for the above state
- The figure also shows the values of all successors, top successors have $h = 12$

Hill-climbing search: 8-queens problem



- A local minimum with $h = 1$ (obtained in 5 steps)

Hill-climbing search

- The algorithm gets stuck for several reasons:
 - Local Maximum: it is a peak that is higher than each of its neighbours, but lower than the maximum overall
 - Ridges: cause a sequence of local maxima that make navigation difficult
 - Plateau (flat): can lead to a local maximum where there is no ascendant exit or a terrace to advance
- In the 8-queens, it gets stuck in 86% and solve 14% cases
- If we allow lateral movements with the hope that we find a terrace (limiting them if reach a local maximum, e.g.100): → 94% success
- Variants:
 - Stochastically (randomly chooses upward movements)
 - Random restart (the initial states are generated randomly)


Informed search strategies

- Best-first search
- Greedy best-first search
- A* search
- Heuristics
- Local search algorithms
 - Hill-climbing search
 - **Simulated annealing search**
 - Local beam search

Simulated annealing search

- Process of tempering or hardening metals by heating and then cooling them gradually
- Idea: escape local maxima by allowing some "bad" moves but **gradually decrease** their frequency
- It combines hill-climbing with random generation successor
- Good for problems with a large search space, optimum is surrounded by many local optima
- Problem: determine the values of the parameters, and requires an important work of experimentation that depends on each problem
- Widely used in VLSI layout, airline scheduling, etc

Simulated annealing search



Informed search strategies

- Best-first search
- Greedy best-first search
- A^* search
- Heuristics
- Local search algorithms
 - Hill-climbing search
 - Simulated annealing search
 - **Local beam search**

Local beam search

- Idea: Keep track of k states rather than just one
 - Start with k randomly generated states
 - At each iteration, all the successors of all k states are generated
 - If any one is a goal state, stop; else select the k best successors from the complete list and repeat
 - Alternatively stochastic LBS randomly choose k successors, with the probability of choosing a successor as an increasing function of its value

Outline

- Introduction
- Problem formulation
- Problem types
- Basic search algorithms
- **Conclusions**

Conclusions

- Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored
- Variety of uninformed and informed search strategies